

# American Fuzzy Lop

Introduction & technical demo



**What is “fuzzy testing” ?**

# What is “fuzzy testing” ?

- **A way to test software products**
  - test sureness : try to find out lot of crash cases that you have not imagined yet
  - test security : in case of crash, be sure user cannot use your soft to break some things around



# What is “fuzzy testing” ?

- **The manner :**
  - providing unexpected, invalid or random data
  - usual inputs : keyboard, mouse, API calls
  - unusual input but think to it : databases, SHM and all other ways that I didn't think of yet



# Two types of fuzzing

- Generation based
- Mutation based



# American Fuzzy Lop

# AFL in few words

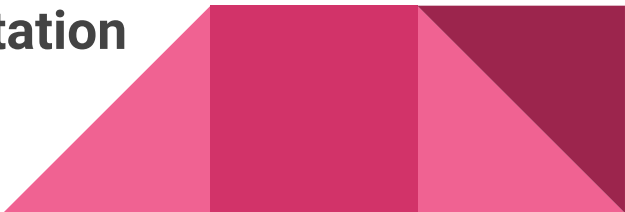
**Mutation based fuzzer**

**Genetic algorithm**

**High efficiency**

**Code instrumentation**

note : every “cf.” notifications, refer to the documentation of afl given with the package afl-2.10b from [this link](#).

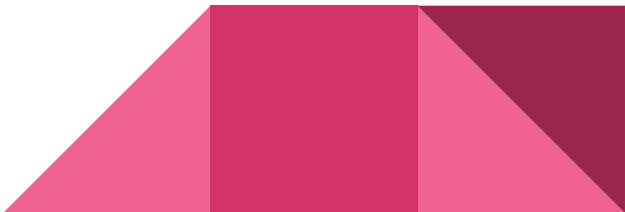


# In which case can it be used ?

## For white box usage

- 1st step : compile your soft with AFL's compilers (based on gcc or clang)
- 2nd step : launch your test
- 3rd step : explore found crash cases

## For black box usage

- instrumentation possible under QEMU
  - 2-5x slower than classic instrumentation
  - some issues with parallelizism
- 



# Compilation

## 2 levels of code instrumentation for target binary

**classic : instrumented at assembly level** (near native execution speed)

afl-gcc            &    afl-g++

afl-clang        &    afl-clang++

**fast : clang only, instrumented at compiler level** (+/-10% better than native speed)

afl-clang-fast    ->should replace afl-clang someday

afl-clang-fast++



# What's needed to run some tests ?

## Must do

- set CPU **scale governor** to “**performance**” (won't work if not)

## Must provide

- One or few ***valid*** input case file(s)

## Keep it simple

- More ***various*** and ***shortest files***, are better to get good perfs (in term of test efficiency)



# Execution

## Command line : level 1

**basic** : give **input files** on **stdin**

```
afl-fuzz -i valid_inputs_folder/ -o afl_output_folder/ ./binary_to_test
```

**AFL** simply **send inputs** on **stdin** of your binary



# Execution

## Command line : level 1

**basic** : give **input files** as **parameter**

```
afl-fuzz -i valid_inputs_folder/ -o afl_output_folder/ ./binary_to_test @@
```

@@ symbolize the **place** of the **input file** in your command



# Execution

## Command line : level 2

Some **options** :

- x a\_dictionnary : put a dictionnary to help fuzzer
- Z 1 : force fuzzer to use specified processor core  
(little performance gain)
- T 50 : set timeout to : 5x reference time + 50 ms  
(basic timeout is : 5x reference time + 20 ms)

Sample of JINK dictionnary

```
full_basic_jink.dict
"tags"
"name"
"segment"
"data"
"label"
"format"
"formats"
"channels"
"metadata"
"type"
"sample-rate"
"uniform-sampling"
"unit"
"resolution"
"quantization"
"min"
```

# Execution

## Command line : level 2


For **multi thread** test :

-**M** a\_name : specify ***master*** fuzzer (deterministic)

-**S** an\_other\_name : specify ***slave*** fuzzer (random)

**Only one master** and as many slaves as you want

**All fuzzers** must **share** the same **input\_folder** as well as the **output\_folder**.



# And then, pull the trigger...

```
bulld@sdkt048-jessie-fuzzy 11:18:59 test2 $ ./test2.sh hey
```

(my script just want name as param  
do not pay attention)

```
af1-fuzz 2.10b by <lcantuf@google.com>
```

-Z 1

```
[+] You have 12 CPU cores and 4 runnable tasks (utilization: 33%).
```

```
[+] Try parallel jobs - see docs/parallel_fuzzing.txt.
```

```
[+] Using specified CPU affinity: main = 1, child = 1
```

```
[*] Checking core_pattern...
```

```
[*] Checking CPU scaling governor...
```

```
[*] Setting up output directories...
```

```
[+] Output directory exists but deemed OK to reuse.
```

```
[*] Deleting old session data...
```

```
[+] Output dir cleanup successful.
```

```
[*] Scanning 'in'...
```

```
[+] No auto-generated dictionary tokens to reuse.
```

```
[*] Creating hard links for all input files...
```

```
[*] Validating target binary...
```

```
[*] Attempting dry run with 'id:000000,orig:myInk.jink'...
```

```
[*] Spinning up the fork server...
```

```
[+] All right - fork server is up.
```

```
len = 518, map size = 6095, exec speed = 3313 us
```

```
[+] All test cases processed.
```

```
[+] Here are some useful stats:
```

```
Test case count : 1 favored, 0 variable, 1 total
```

```
Bitmap range : 6095 to 6095 bits (average: 6095.00 bits)
```

```
Exec timing : 3313 to 3313 us (average: 3313 us)
```

no -x given

test binary with input cases

what is written

```
[*] No -t option specified, so I'll use exec timeout of 20 ms.
```

```
[+] All set and ready to roll!
```

Everything is ok !

# And then, pull the trigger...

```
american fuzzy lop 2.10b (more_heterogeneous_jink)

process timing
  run time : 0 days, 17 hrs, 45 min, 22 sec
  last new path : 0 days, 0 hrs, 27 min, 23 sec
  last uniq crash : 0 days, 0 hrs, 6 min, 54 sec
  last uniq hang : 0 days, 5 hrs, 28 min, 54 sec
cycle progress
  now processing : 2042 (98.55%)
  paths timed out : 0 (0.00%)
stage progress
  now trying : interest 32/8
  stage execs : 133k/198k (67.07%)
total execs : 134M
(exec speed : 332.5/sec)
fuzzing strategy yields
  bit flips : 600/3.27M, 75/3.27M, 16/3.27M
  byte flips : 0/408k, 2/408k, 5/407k
  arithmetics : 117/22.7M, 2/2.31M, 0/4549
  known ints : 21/1.93M, 18/11.2M, 12/17.7M
  dictionary : 95/22.1M, 46/22.3M, 48/20.1M
  havoc : 120/2.87M, 0/0
  trim : 34.22%/167k, 0.09%

overall results
  cycles done : 0
total paths : 2072
  uniq crashes : 90
  uniq hangs : 40

map coverage
  map density : 9497 (14.49%)
  count coverage : 2.62 bits/tuple

findings in depth
  favored paths : 328 (15.83%)
  new edges on : 548 (26.45%)
total crashes : 297k (90 unique)
total hangs : 1558 (40 unique)

path geometry
  levels : 6
  pending : 1683
  pend fav : 16
  own finds : 1087
  imported : 984
  variable : n/a

[cpu@00: 19%]
```

**total path** : unique call stacks found

**total execs** : number of tests done

**“90 unique”** : means 90 different call stacks that crash

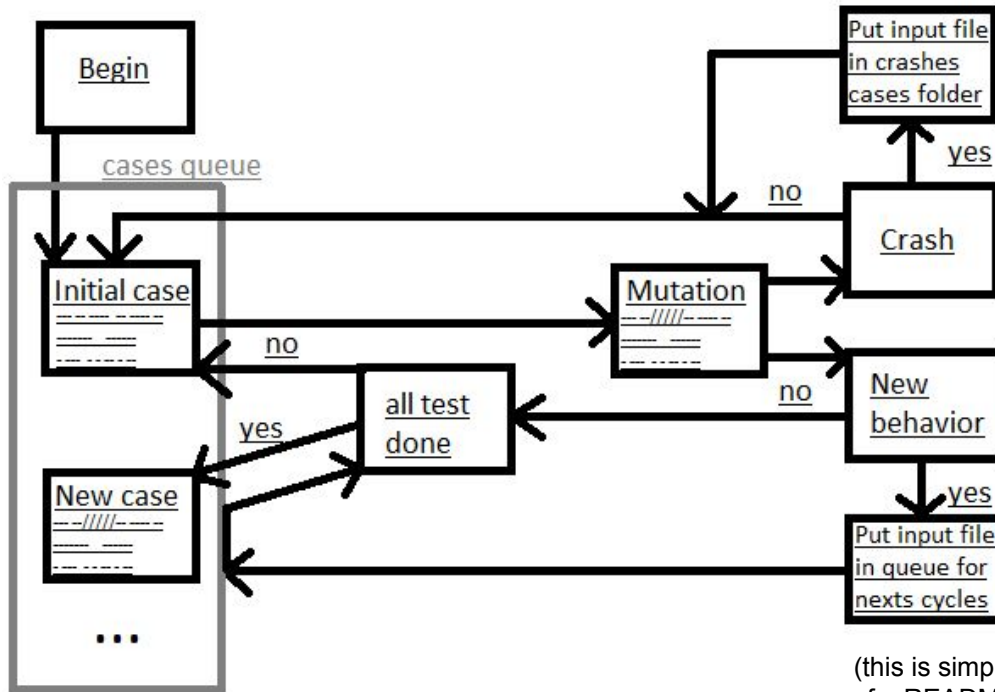
**levels** : deepness of the mutation

**exec speed** : number of tests / sec



# Nice ! It works :)

## But... what's really happening ?



(this is simplistically draw)  
cf. : README

# Mutations ? What's mutation ?

- **bitflip L/S** : L bits are flipped every S bits, variants are 1/1, 2/1, 4/1, 8/8, 16/8, 32/8
- **arith L/8** : try to subtract or add smalls integer to 8, 16 or 32 bits
- **interest L/8** : same as arith but with “interesting” values
- **extras** : in case of given dictionnary it will overwrite or insert values
- **havoc** : fixed length cycle that randomly combine every precedent technics, and add : bloc deletion and/or bloc duplication
- **splice** : last resort strategy, in case that no new path were found for an entire cycle : do the same as havoc, but first splices together two random inputs from the queue at some arbitrarily selected midpoint

cf. : status\_screen.txt



# One more thing ...

**Software with heavy initial load can get huge performance boost**  
with :

## **AFL persistent mode**

For me, it make the fuzzer increase by **1000%** the average speed execution.

**From 300 to 3000...**

It's easy to set up, just go read [README.llvm](#).



# Lot of other things

There is a lot of other tools & functionality provided by AFL, i can't list them all.

There is a lot a way to use them for many reasons.

Just go to <http://lcamtuf.coredump.cx/afl/> get the last release and read all the doc available.

Also, be part of the AFL project by asking questions and proposing ideas for this project, on the google group next here :

<https://groups.google.com/forum/#!forum/afl-users>



# Thank you for reading

I hope you learnt some interesting things.

This is a completely subjective documentation.  
It's absolutly not the official one or the unique. And it's no more the only way  
to use AFL.

HOUDU Loïc  
loic.houdu@gmail.com

